

PII: Timing a Linux Process via Socket

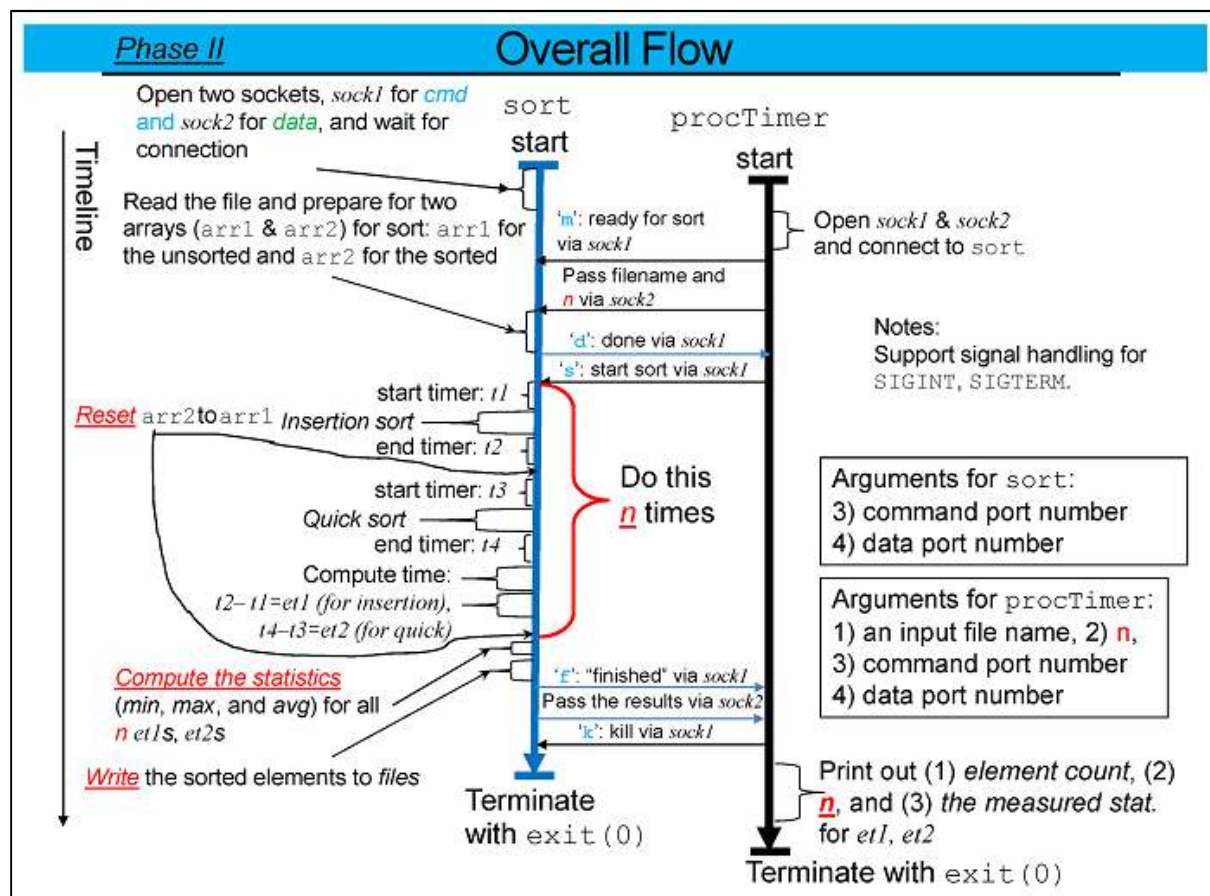
ELEC462004 System Programming – Term Project

(Instructor: Prof. Suh, Young-Kyoon)

Project Phase II: Due:11:59pm (midnight) December 6, 2017

In Phase II, you are to implement the timing program on measuring the elapsed time (ET) (Phase I) of a portion of running an algorithm via SOCKET. (Note that the previous phase II that was described before is replaced by this.) Your project must be implemented in C and be running on Linux (preferably Ubuntu).

Task description



- Write a C program, named `sort`, to perform insertion and quick sorting tasks on a given datafile. It takes two arguments: one (*port1*) for command and another (*port2*) for data. It then opens two sockets: *sock1* at *port1* and *sock2* at *port2*. It then waits for the connection from the below program.

[Usage] `./sort port_number_for_command port_number_for_data`

Ex. `./sort 2400 2401`

- Write another C program, named `procTimer`, which prints out the *minimum*, *maximum*, and *average* execution time of a sorting task “in memory” (done in another program in the above). It takes four arguments: one for the name of a datafile for the sorting, another for how many times the sorting should be performed (*n*), and the other two for command and data ports, `port1` and `port2`. Note that *port1* and *port2* should be the same as those two from `sort`. Below are the expected command-line arguments.

```
[Usage] ./procTimer          input_data_file          number_of_repetitions
port_number_for_command port_number_for_data
```

```
Ex. ./procTimer sample.txt 1000 2400 2401
```

`procTimer` opens two sockets: *sock1* at *port1* and *sock2* at *port2*. It then connects to `sort` through both sockets. `procTimer` sends one character of ‘m’ to `sort` via *sock1* to make it ready for its forthcoming sorting task. It subsequently sends (i) the name of the given input data file (e.g. `sample.txt`) and (ii) the number of repetitions (e.g. *n*) to `sort` via *sock2*.

- `sort` reads the input data file (named `sample.txt`) and then fills all the elements in an unsorted data array (named `arr1`). It also creates another array (named `arr2`) to contain the elements to be sorted. `sort` initializes `arr2` by copying all the elements in `arr1` by the *reset* operation. It then sends one character of ‘d’ to `procTimer` via *sock1* to say that it’s ready for `sort`.
- `procTimer` sends one character of ‘s’ to `sort` via *sock1* to tell it to “start” the sorting.
- Finish the sorting at `sort` as done in Phase I. Make sure that you should reset `arr2` to `arr1` after each repetition for each `sort`.
- Once `sort` finishes the measurement, it computes the statistics (the minimum, maximum, and average) for each `sort`. `sort` then writes the sorted elements into the files, named, `insertion_sort_res.txt` for insertion sort, and another, named `quick_sort_res.txt` for quick sort. Both files must have the same sorted data files: use ‘diff’ to make certain that both have no difference.
- Finally, `sort` sends one character of ‘f’ to `procTimer` via *sock1* to say that the sorting is “finished.” It then sends the measured data to `procTimer` via *sock2*.
- `procTimer` shows on screen the receive data through standard out. Finally, `procTimer` sends one character of ‘k’ to `sort` via *sock1* to kill it and then it terminates with an exit value of 0. Then `sort` terminates with an exit value of 0 on receipt of the character (‘k’) sent from `procTimer`.

- The output looks like the following:

```
// one terminal

~$./sort 2400 2401 (or '&' can be appended too.)

(Now it waits...)

Received 'm'

Received 's'

Received 'k'

(It dies now.)

~$
```

```
// another terminal

~$./procTimer sorting_input.txt 10000 2400 2401

Received 'd'

Received 'f'

### Received the measurement statistics ###

The elements in the input data file: 100000

The number of repetitions: 10000

Insertion sort

- The minimum elapsed time: 10000 milliseconds
- The maximum elapsed time: 11000 milliseconds
- The average elapsed time: 11000 milliseconds

Quick sort

- The minimum elapsed time: 9000 milliseconds
- The maximum elapsed time: 10000 milliseconds
- The average elapsed time: 9500 milliseconds

~$
```

Notes

- Your code MUST use dynamic allocation. We'll check memory leak after running your code.
- Again, `sort` and `procTimer` must support signal handling on the following signals: `SIGINT` and `SIGTERM`.
 - o Those two signals must be ignored.
- We provide you with an input file having 100,000 integers. See the `project_phase2` directory.
- [**IMPORTANT!!**] You must include into your submission directory (to be described shortly) your **Makefile** to produce `sort` and `procTimer`.
 - o If you don't have your make file, then you get **ZERO credit** for this phase 2.

<WARNING!!>

- Suppose that your leader's studentID is '**2014123123**' and the home directory for phase 2: `PROJP2_HOME=/usr/share/fall17/elec462004/project_phase2`.
- Create your directory named the leader's student ID under `PROJP2_HOME`: specifically, the name of the submission directory will be `$PROJP2_HOME/2014123123` in the above example.
- To complete your submission, put all relevant project code into the directory that you just created above, which will be like `$PROJP2_HOME/2014123123` in the above example.
- Make sure you have your **Makefile** to generate the executables: `sort` and `procTimer`.
- If you use the wrong name of the directory or place your files in the wrong directory,
NO grade will be given.
- Since this is in-memory sorting, we'll run your code against a **HUGE VOLUME** of elements (at least 50K elements) stored in a data file.
- - Before you finish all the dynamic data structures, you should free them. We will check if you have any **memory leak** at the end of the program via `valgrind`. We'll give you a penalty of **5%** if the leak is detected.

QnA

Feel free to ask whatever questions you may have to me by email. Or you could visit my office by appointment.

Late Day Policy

All exercises are due at midnight on the scheduled due date. A grading penalty will be applied to late assignments. Any assignment turned in late will be penalized 50% per late day.

Plagiarism

No plagiarism will be tolerated. If the assignment is to be worked on your team own, please respect it. If the instructor determines that there are substantial similarities exceeding the likelihood of such an event, he will call the two (or more) students to explain them and possibly to take an immediate test (or assignment, at the discretion of the instructor) to determine the student's abilities related to the offending work.